

DATA() {  
EXPOSED;

# ADF Data Flows Best Practices and Performance Optimizations

Mark Kromer @kromerbigdata

*ADF Program Manager*



# Agenda

- Azure SQL DB ETL Performance
- Transformation optimizations
- Monitoring
- Global Settings
- Best Practices
- Azure Integration Runtimes

# Database ETL Performance

# Sample timings for Azure SQL DB

## Scenario w/Azure SQL DB

- Source: Azure SQL DB Table
- Sink: Azure SQL DB Table
- Table size: 74 columns, 887k rows
- Transforms: Single derived column to mask 3 fields
- Time: 3 mins end-to-end using memory optimized 80-core debug Azure IR
- Recommended settings: Source partitioning on SQL DB Source, current partitioning on Derived Column and Sink

# SQL Database Timing

## Performance run for csv to Azure sql

Dataset size			15GB		Dataset size			31.44 GB	
Storage type			csv to Azsql		Storage type			Azsql to csv	
Total Rows			71 Million		Total Rows			71 Million	
Total Partition			114		Total Partition			6,8	
Compute type			General Purpose		Compute type			General Purpose	
Sink type			Azsql		Sink type			CSV	
AZSQL Type			Genral purpose: Gen5, 16 vCores		AZSQL Type			Genral purpose: Gen5, 16 vCores	
					azsql to csv with optimize source partition(6)			azsql to csv with optimize source partition(8)	
Core		Stage time in Minutes		Sink time in Minutes		sink time in minutes		Sink time in minutes	
8		70.47		94.15		9.41		10.34	
16		42.12		69		9.25		8.18	
32		29.29		55.25		8.3		6.49	
64		22.7		44		8.49		7.15	

**Note:**Cluster uptime not included

**Sink time is total job execution time**

Total job exection time is only sink time

After selecting the optimize with source 6 partition performace has improved

# Synapse DW Timing

Dataset size	15GB		45GB		Dataset size	15GB
Storage type	CSV TO DW		Blob to Blob		Storage type	DW TO CSV
Total Rows	71 Million		337 million		Total Rows	71 Million
Performance level	Gen2: DW1000c		Gen2: DW1000c		Performance level	Gen2: DW1000c
Total partition	114		358		Total partition	1
Core	sink stage time in minutes	sink time in minutes	Sink stage time in minutes	sink time in minutes	stage time in minutes	sink time in minutes
8	5.11	30.54	24.23	90.41	45.57	46
16	3.18	30.1	13.59	79.17	42.27	42.28
32	3.23	27.2	7.36	81.16	42.11	42.12
64	1.13	26.12	5.14	77.55	45.1	45.13

Compute type: General Purpose

Note:-

**Cluster uptime not included**

**Sink time is total job execution time**

Adding cores proportionally decreases time it takes to process data into staging files for Polybase. However, there is a fairly static amount time that it takes to write that data from Parquet into SQL tables using Polybase.

# CosmosDB Timing

## Performance run for csv to Cosmos db and json to Cosmos DB

Dataset size	15GB	45GB(csv data converted to json )	Dataset size	15GB
Storage type	CSV TO COSMOS	JSON TO COSMOS (multi region write enabled)	Storage type	COSMOS TO CSV
Total Rows	71 Million	71 Million	Total Rows	71 Million
Cosmos DB Throughput	100000	100000	Cosmos DB Throughput	100000

Core	sink time in minutes	Number of Partition	Number of partition	Sink time in minutes	Number of partition	Sink time in minutes
8	300	114	352	448	10	34.56
16	280	114	352	316	10	21.37
32	270	114	352	257	10	18.28
64	260	114	352	238	10	18.37
128	301	128	352	226		
256	400	255	352	225		

Note:- Out put data genrated in cosmos db:170 GB      Data size:123:52 GB      Index size:45:31 GB

**Cluster uptime not included**

**Sink time is total job execution time**

**Note:- when input dataset in csv format if number of expected partition is less then cpu core then performance goes down due to more number of partition data got created**

Note:- From csv to cosmos DB with multi region write enabled and disable results are same

when cpu core is more then number of partition performance get decreased

**Compute type: General Purpose**

# Transformation Performance



# Window / Aggregate Timing

Dataset size	76GB	
Storage type	ADLSGEN2	
Total Rows	600Million	600013353
Result time	Sink time	
Compute type	General Purpose	
Sink type	CSV	
Total Partition	199	197


  

Core	window	aggregate+join
8	43.44	54
16	23.11	27.17
32	12.49	14.36
64	7.40	8.51

Tpch dataset has been used for tests  
Lineitem table size 76Gb  
Aggregate on sum L\_TAX  
window on L\_SUPPLYKEY

## Compute type: General Purpose

- Performance improvement scales proportionately with increase in Vcores
- 8 Vcore to 64 Vcore performance increase is around 5 times more



# Transformation Timings

## Performance run for CSV Transformation

Dataset size	15GB,20GB
Storage type	ADLSGEN2
Total Rows	71 Million
Result time	Sink time
Total Partition	200

Core	csv left join	csv lookup	Exists
8	14.26	28.49	12.5
16	7	22.45	12.48
32	5.3	9.5	5.15
64	2.44	5.38	3.14

Note: - Each run has ran sequentially

**Cluster uptime not included**

Note: Time in Minutes, All number is avrage of two runs

Goal: Find which transformation is faster among these left join, lookup, Exists. These all transformation has same output

## Compute type: General Purpose

Dataset size	15GB,20GB
Storage type	ADLSGEN2
Total Rows	71 Million
Result time	Sink time
Total Partition	200

Core	window	aggregate+join (two source)	aggregate + join (one source)
8	8.25	10.49	10.19
16	7.56	7.1	6.25
32	3.4	3.54	4.12
64	2.44	2.46	2.21

## Transformation recommendations


- When ranking data across entire dataset, use Rank transformation instead of Window with rank()
- When using rowNumber() in Window to uniquely add a row counter to each row across entire dataset, instead use the Surrogate Key transformation

# TPCH Timings

## TPCH CSV in ADLS Gen 2

Dataset size	76GB,16GB		
Storage type	ADLS GEN2		
Total Rows	600Million	600013353	
Result time	Sink time		
Total Partition	175	190	200
			<div>Tpch dataset has been used for tests Lineitem table size 76Gb Order Table size 16Gb Column O_ORDERKEY</div>
Core	leftouter join	Exists	lookup
8	52.23	41.9	74.57
16	41.59	26.55	45.29
32	33.53	20.49	26.53
64	29.59	18.23	19.17

## Compute type: General Purpose



# Optimizing transformations

## **Each transformation has its own optimize tab**

- Generally better to not alter -> reshuffling is a relatively slow process

## **Reshuffling can occur if data is very skewed**

- One node has a disproportionate amount of data

## **For Joins, Exists and Lookups:**

- If you have a many of these transforms, memory optimized greatly increases performance
- Use cached lookup w/cached sink
- Can 'Broadcast' if the data on one side is small
- Rule of thumb: Less than 50k rows

## **Use Window transformation partitioned over segments of data**

- For Rank() across entire dataset, use the Rank transformation instead
- For RowNumber() across entire dataset, use the Surrogate Key transformation instead

## **Transformations that require reshuffling like Sort negatively impact performance**

# Azure Data Factory Data Flow Performance

*\* Includes cold cluster start-up time*

IR Size	ADLS Source File				Load						Sink							
	File Type	Single File Size	File Count	Total MB	Type	Partition		Load Time		MB/s	Type	Partition		sink process time		MB/s		
						num	size	min	sec			num	size	min	sec			
8-8-general	small	1MB	51200	50807	CDM	1970	26MB	12	4	70	CDM	1970	26MB	14	8	60		
				50807	CDM	1970	26MB	11	28	74	Dataset	1970	26MB	13	22	63		
				50807	Dataset	1970	26MB	11	36	73	CDM	1970	26MB	12	51	66		
				50807	Dataset	1970	26MB	9	28	89	Dataset	1970	26MB	10	44	79		
8-8-general	medium	242MB	210	50791	CDM	210	242MB	6	37	128	CDM	210	242MB	6	47	125		
				50791	CDM	210	242MB	6	44	126	Dataset	210	242MB	6	55	122		
				50791	Dataset	210	242MB	7	22	115	CDM	210	242MB	7	40	110		
				50791	Dataset	210	242MB	5	50	145	Dataset	210	242MB	6	10	137		
8-8-general	large	1067MB	50	50794	CDM	50	1067MB	6	23	133	CDM	50	1067MB	6	27	131		
				50794	CDM	50	1067MB	7	6	119	Dataset	50	1067MB	7	10	118		
				50794	Dataset	50	1067MB	8	6	105	CDM	50	1067MB	8	23	101		
				50794	Dataset	50	1067MB	5	31	153	Dataset	50	1067MB	5	49	146		
				50807	CDM	1970	26MB	5	47	146	CDM	1970	26MB	7	59	106		
				50807	CDM	1970	26MB	5	49	146	Dataset	1970	26MB	7	49	108		
				50807	Dataset	1970	26MB	5	40	149	CDM	1970	26MB	7	3	120		
				50807	Dataset	1970	26MB	4	21	195	Dataset	1970	26MB	6	13	136		
16-16-general	medium	242MB	210	50791	CDM	210	242MB	3	45	226	CDM	210	242MB	3	54	217		
				50791	CDM	210	242MB	3	50	221	Dataset	210	242MB	4	9	204		
				50791	Dataset	210	242MB	3	58	213	CDM	210	242MB	4	16	198		
				50791	Dataset	210	242MB	2	42	314	Dataset	210	242MB	3	9	269		
16-16-general	large	1067MB	50	50794	CDM	50	1067MB	4	19	196	CDM	50	1067MB	4	22	194		
				50794	CDM	50	1067MB	3	54	217	Dataset	50	1067MB	3	59	213		
				50794	Dataset	50	1067MB	4	24	192	CDM	50	1067MB	4	41	181		
				50794	Dataset	50	1067MB	2	43	312	Dataset	50	1067MB	3	2	279		
4-4-general	small	1MB	51200	50807	CDM	1970	26MB	21	37	39	CDM	1970	26MB	23	31	36		
				50807	CDM	1970	26MB	21	8	40	Dataset	1970	26MB	22	54	37		
				50807	Dataset	1970	26MB	19	24	44	CDM	1970	26MB	20	48	41		
				50807	Dataset	1970	26MB	15	5	56	Dataset	1970	26MB	16	19	52		
4-4-general	medium	242MB	210	50791	CDM	210	242MB	12	17	69	CDM	210	242MB	12	27	68		
				50791	CDM	210	242MB	12	31	68	Dataset	210	242MB	12	50	66		
				50791	Dataset	210	242MB	12	48	66	CDM	210	242MB	13	5	65		
				50791	Dataset	210	242MB	9	24	90	Dataset	210	242MB	9	50	86		
4-4-general	large	1067MB	50	50794	CDM	50	1067MB	12	45	66	CDM	50	1067MB	12	48	66		
				50794	CDM	50	1067MB	12	34	67	Dataset	50	1067MB	12	50	66		
				50794	Dataset	50	1067MB	12	58	65	CDM	50	1067MB	13	15	64		
				50794	Dataset	50	1067MB	9	6	93	Dataset	50	1067MB	9	24	90		

# Azure Synapse Data Flow Performance

*\* Includes cold cluster start-up time*

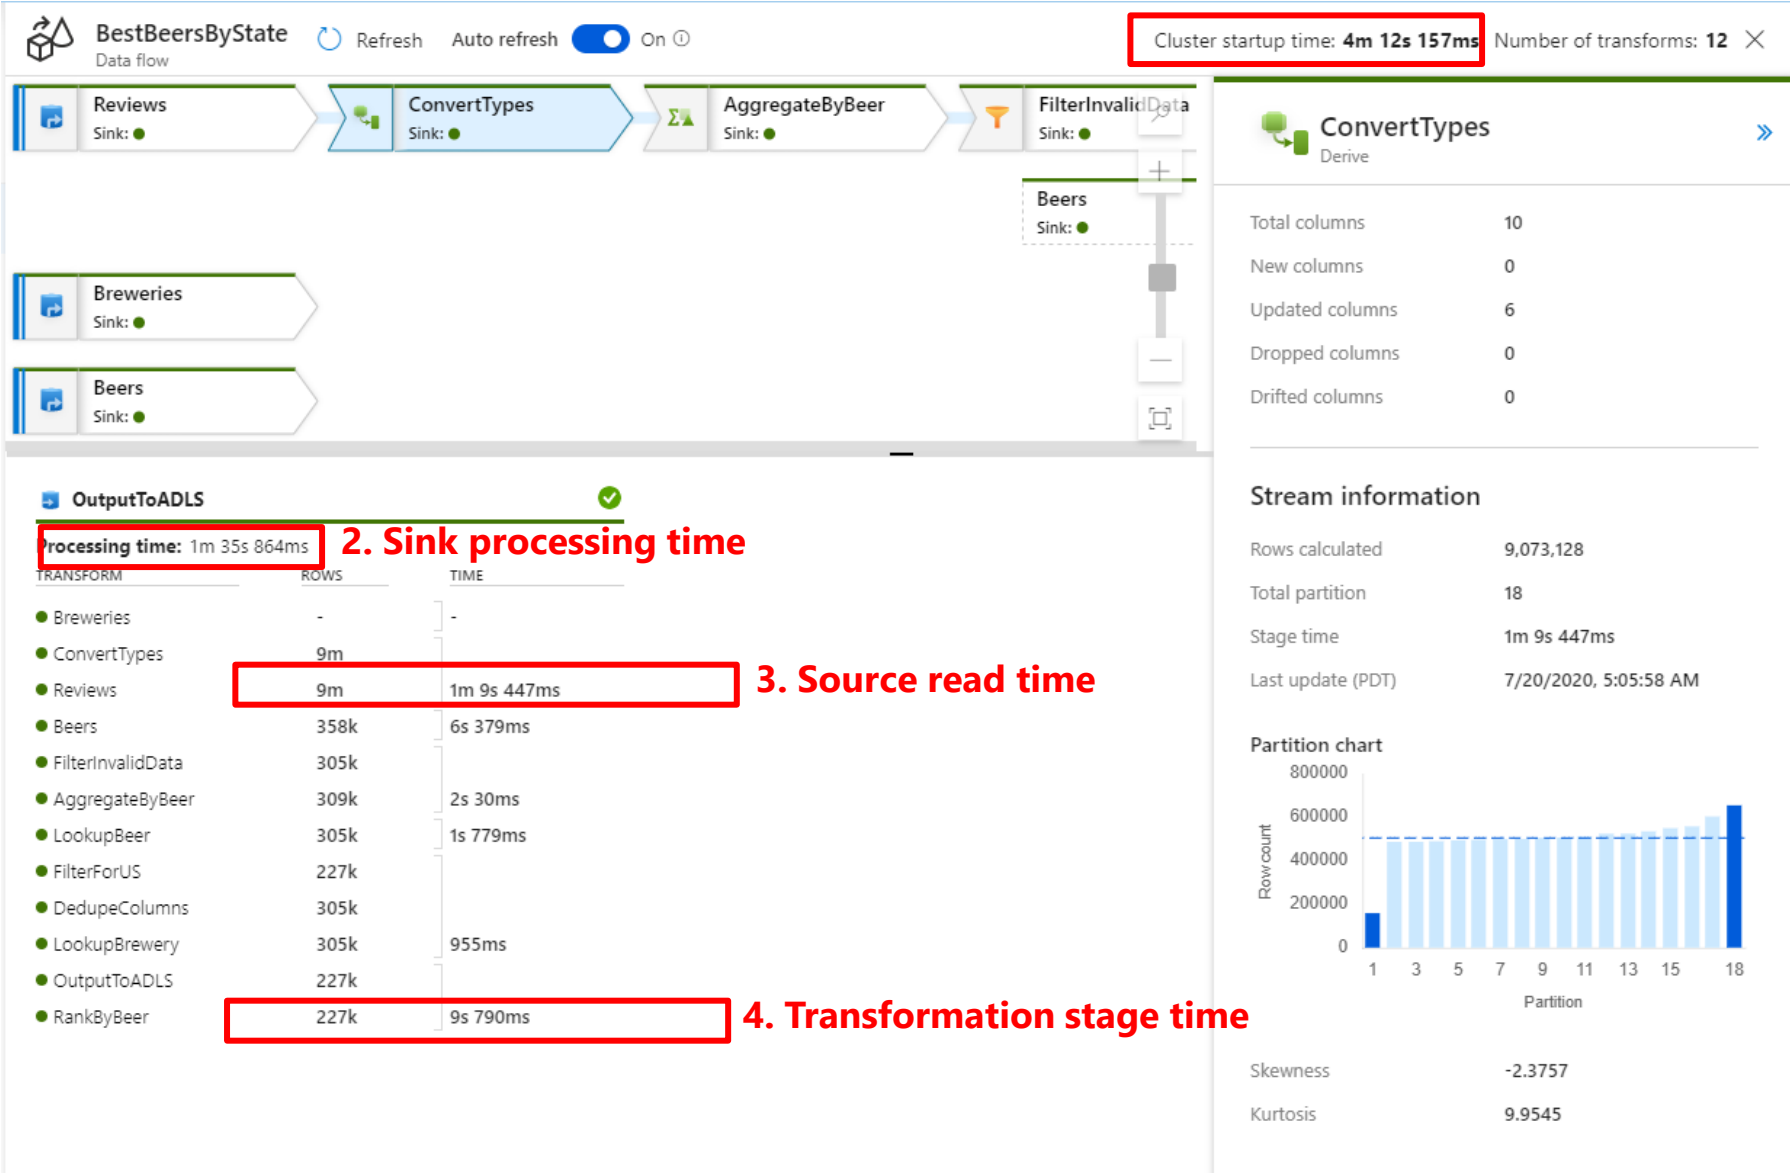
IR Size	ADLS Source File				Load						Sink					
	File Type	Single File Size	File Count	Total MB	Type	Partition		Load Time		MB/s	Type	Partition		sink process time		MB/s
						num	size MB	min	sec			num	size MB	min	sec	
8-8-general	small	1MB	51200	50807	CDM	250	204	10	36	80	CDM	250	204	11	20	75
				50807	CDM	250	204	11	21	75	Dataset	250	204	12	1	70
				50807	Dataset	250	204	12	9	70	CDM	250	204	12	19	69
				50807	Dataset	250	204	7	56	107	Dataset	250	204	8	8	104
8-8-general	medium	242MB	210	50791	CDM	53	958	8	47	96	CDM	53	958	8	51	96
				50791	CDM	53	958	8	15	103	Dataset	53	958	8	20	102
				50791	Dataset	53	958	8	55	95	CDM	53	958	8	58	94
				50791	Dataset	53	958	5	31	153	Dataset	53	958	5	35	152
8-8-general	large	1067MB	50	50794	CDM	50	1067	9	25	90	CDM	50	1067	9	28	89
				50794	CDM	50	1067	8	44	97	Dataset	50	1067	8	48	96
				50794	Dataset	50	1067	8	39	98	CDM	50	1067	8	42	97
				50794	Dataset	50	1067	6	8	138	Dataset	50	1067	6	11	137
16-16-general	small	1MB	51200	50807	CDM	250	204	5	54	144	CDM	250	204	6	30	130
				50807	CDM	250	204	5	52	144	Dataset	250	204	6	30	130
				50807	Dataset	250	204	6	27	131	CDM	250	204	6	38	128
				50807	Dataset	250	204	4	21	195	Dataset	250	204	4	30	188
16-16-general	medium	242MB	210	50791	CDM	53	958	5	26	156	CDM	53	958	5	30	154
				50791	CDM	53	958	5	21	158	Dataset	53	958	5	26	156
				50791	Dataset	53	958	5	22	158	CDM	53	958	5	25	156
				50791	Dataset	53	958	5	40	149	Dataset	53	958	5	44	148
16-16-general	large	1067MB	50	50794	CDM	50	1067	5	45	147	CDM	50	1067	5	49	146
				50794	CDM	50	1067	5	15	161	Dataset	50	1067	5	19	159
				50794	Dataset	50	1067	5	32	153	CDM	50	1067	5	35	152
				50794	Dataset	50	1067	4	24	192	Dataset	50	1067	4	27	190
4-4-general	small	1MB	51200	50807	CDM	250	204	20	36	41	CDM	250	204	21	33	39
				50807	CDM	250	204	24	57	34	Dataset	250	204	25	55	33
				50807	Dataset	250	204	23	29	36	CDM	250	204	23	41	36
				50807	Dataset	250	204	16	20	52	Dataset	250	204	16	38	51
4-4-general	medium	242MB	210	50791	CDM	53	958	17	58	47	CDM	53	958	18	3	47
				50791	CDM	53	958	17	34	48	Dataset	53	958	17	41	48
				50791	Dataset	53	958	19	36	43	CDM	53	958	19	39	43
				50791	Dataset	53	958	10	41	79	Dataset	53	958	10	45	79
4-4-general	large	1067MB	50	50794	CDM	50	1067	17	21	49	CDM	50	1067	17	27	49
				50794	CDM	50	1067	18	46	45	Dataset	50	1067	18	51	45
				50794	Dataset	50	1067	19	4	44	CDM	50	1067	19	8	44
				50794	Dataset	50	1067	11	0	77	Dataset	50	1067	11	3	77

# ETL Performance Monitoring

# Identifying bottlenecks

## 1. Cluster startup time

1. Sequential executions can lower the cluster startup time by setting a TTL in Azure IR
2. Total time to process the stream from source to sink. There is also a post-processing time when you click on the Sink that will show you how much time Spark had to spend with partition and job clean-up. Write to single file and slow database connections will increase this time
3. Shows you how long it took to read data from source. Optimize with different source partition strategies
4. This will show you bottlenecks in your transformation logic. With larger general purpose and mem optimized IRs, most of these operations occur in memory in data frames and are usually the fastest operations in your data flow





# Global configurations that effect performance

## **Logging level (pipeline activity)**

- Verbose (default) is most expensive
- You can get a small increase in performance for large data flows without detailed logging
- Trade-off: Less diagnostics

## **Error row handling (sink transformation)**

- Expect 5%-10% perf hit
- Trade-off: Provides detailed logging and continuation of data flow on database driver errors

## **Run in parallel (pipeline activity)**

- Currently only available for “connected” streams, i.e. multiple sinks from a single stream
- Can write to multiple sinks at same time
- Use with new branch, conditional split

## **Parallel activity executions (pipeline activity)**

- If you place data flow activities on your pipeline canvas without connector lines, your data flows can all start at the same time, lowering overall pipeline execution times.

# ETL Performance Best Practices

## Best practices - Sources

**When reading from file-based sources, data flow automatically partitions the data based on size**

~128 MB per partition, evenly distributed

Use current partitioning will be fastest for file-based and Synapse using PolyBase

Enable staging for Synapse

**For Azure SQL DB, use Source partitioning on column with high cardinality**

Improves performance, but can saturate your source database

**Reading can be limited by the I/O of your source**

# Best practices – Debug (Data Preview)

## Data Preview

Data preview is inside the data flow designer transformation properties

Uses row limits and sampling techniques to preview data from a small size of data

Allows you to build and validate units of logic with samples of data in real time

You have control over the size of the data limits under Debug Settings

If you wish to test with larger datasets, set a larger compute size in the Azure IR when switching on "Debug Mode"

Data Preview is only a snapshot of data in memory from Spark data frames. This feature does not write any data, so the sink drivers are not utilized and not tested in this mode.

# Best practices – Debug (Pipeline Debug)

## Pipeline Debug

Click debug button to test your data flow inside of a pipeline

Default debug limits the execution runtime so you will want to limit data sizes

Sampling can be applied here as well by using the “Enable Sampling” option in each Source

Use the debug button option of “use activity IR” when you wish to use a job execution compute environment

This option is good for debugging with larger datasets. It will not have the same execution timeout limit as the default debug setting

# Best practices - Sinks

## SQL:

Disable indexes on target with pre/post SQL scripts

Increase SQL capacity during pipeline execution

Enable staging when using Synapse

Use Source Partitioning on Source under Optimize

Set number of partitions based on size of IR

## File-based sinks:

Use current partitioning allows Spark to create output

Output to single file is a slow operation

Often unnecessary by whoever is consuming data

Can set naming patterns or use data in column

Any reshuffling of data is slow

## Cosmos DB

Set throughput and batch size to meet performance requirements

Clear the folder	<input type="checkbox"/>
File name option *	<input type="radio"/> Default <input type="radio"/> Pattern <input type="radio"/> Per partition <input type="radio"/> As data in column <input checked="" type="radio"/> Output to single file
Output to single file *	<input type="text" value="Enter file name here"/> ⓘ
Quote All	<input type="checkbox"/> ⓘ

Update method	<input checked="" type="checkbox"/> Allow insert <input type="checkbox"/> Allow delete <input type="checkbox"/> Allow upsert <input type="checkbox"/> Allow update
Table action	<input checked="" type="radio"/> None <input type="radio"/> Recreate table <input type="radio"/> Truncate table
Batch size	<input type="text"/> ⓘ
Pre SQL scripts	<input type="text"/>
Post SQL scripts	<input type="text"/>

# Azure Integration Runtime Best Practices

**Data Flows use JIT compute to minimize running expensive clusters when they are mostly idle**

Generally more economical, but each cluster takes ~4 minutes to spin up

IR specifies what cluster type and core-count to use

Memory optimized is best, compute optimized doesn't generally work for production workloads

**When running Sequential jobs utilize *Time to Live* to reuse cluster between executions**

Keeps compute resources alive for TTL minutes after execution for new job to use

Maximum one job per cluster

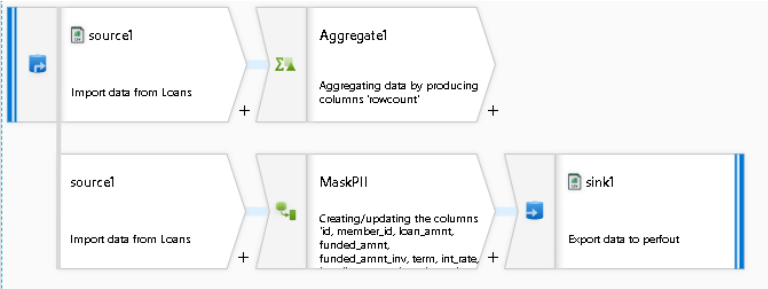
Reduces job startup latency to ~1.5 minutes

**Rule of thumb: start small and scale up**

# Azure IR – General Purpose

- This was General Purpose 4+4, the default auto resolve Azure IR
- For prod workloads, GP is usually sufficient at >= 16 cores
- You get 1 driver and 1 worker node, both with 4 vcores
- Good for debugging, testing, and many production workloads
- Tested with 887k row CSV file with 74 columns
- Default partitioning
  - Spark chose 4 partitions
- Cluster startup time: 4.5 mins
- Sink IO writing: 46s
- Transformation time: 42s
- Sink post-processing time: 45s

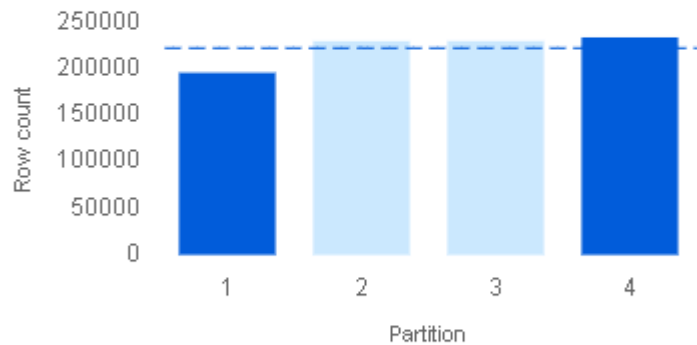
TYPE	RUN START	DURATION	STATUS	INTEGRATION RUNTIME
ExecuteDataFlow	2020-08-11T00:50:36.48789	00:06:17	✔ Succeeded	DefaultIntegrationRuntime (East US)



### Stream information

Rows calculated	887,379
Total partition	4
Stage time	42s 669ms
Last update (PDT)	8/10/2020, 5:56:24 PM

### Partition chart

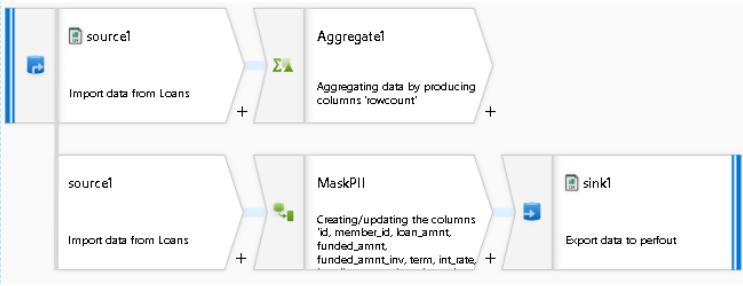


Skewness	-0.9657
Kurtosis	1.7295
Sink processing time	46s 943ms



# Azure IR – Compute Optimized


- Computed Optimized intended for smaller workloads
- 8+8, this is smallest CO option and you get 1 driver and 2 workers
- Not suitable for large production workloads
- Tested with 887k row CSV file with 74 columns
- Default partitioning
  - Spark chose 8 partitions
- Cluster startup time: 4.5 mins
- Sink IO writing: 20s
- Transformation time: 35s
- Sink post-processing time: 40s
- More worker nodes gave us more partitions and better perf than General Purpose



### Stream information

Rows calculated	887,379
Total partition	8
Stage time	23s 901ms
Last update (PDT)	8/10/2020, 6:28:37 PM

### Partition chart

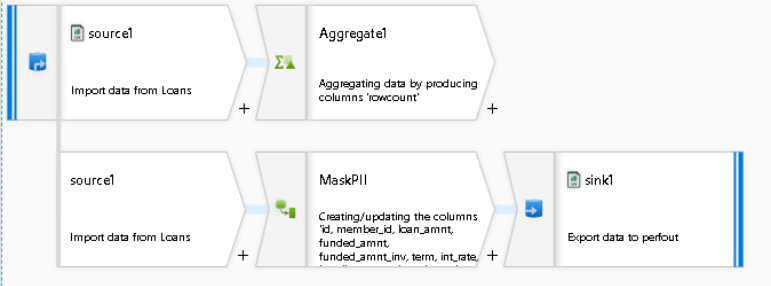


Skewness	-0.9488
Kurtosis	2.6823
Sink processing time	25s 638ms

TYPE	RUN START	DURATION	STATUS	INTEGRATION RUNTIME
ExecuteDataFlow	2020-08-11T01:23:35.76244	00:05:37	✔ Succeeded	DefaultIntegrationRuntime (East US)

# Azure IR – Memory Optimized

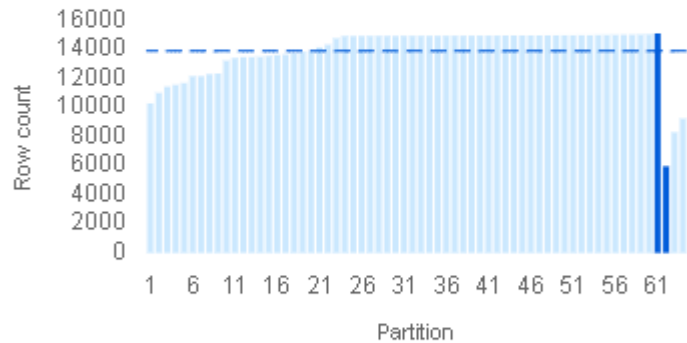
- Memory Optimized well suited for large production workload
- reliability with many aggregates, lookups, and joins
- 64+16 gives you 16 vcores for driver and 64 across worker nodes
- Tested with 887k row CSV file with 74 columns
- Default partitioning
- Spark chose 64 partitions
- Cluster startup time: 4.8 mins
- Sink IO writing: 19s
- Transformation time: 17s
- Sink post-processing time: 40s



## Stream information

Rows calculated	887,379
Total partition	64
Stage time	17s 42ms
Last update (PDT)	8/10/2020, 6:40:58 PM

## Partition chart



Skewness	-2.2327
Kurtosis	8.2938
Sink processing time	19s 254ms

TYPE	RUN START	DURATION	STATUS	INTEGRATION RUNTIME
ExecuteDataFlow	2020-08-11T01:35:20.95875	00:06:09	✔ Succeeded	DefaultIntegrationRuntime (East US)

# Resources

## Complete Data Flows Performance Tuning and Profiles Deck

<https://www2.slideshare.net/kromerm/azure-data-factory-data-flow-performance-tuning-101>

## Data Flows Training

<https://www2.slideshare.net/kromerm/azure-data-factory-data-flows-training-sept-2020-update>

## Data Flows Video Tutorials

<https://docs.microsoft.com/en-us/azure/data-factory/data-flow-tutorials>

## Data Flows Performance Home Page

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-data-flow-performance>

## Copy Data Performance Guidance

<https://docs.microsoft.com/en-us/azure/data-factory/copy-activity-performance>

DATA() {  
EXPOSED;

Learn with us!

View our on-demand playlist:  
[aka.ms/azuresqlandadf](https://aka.ms/azuresqlandadf)

@AzureSQL  
@AzDataFactory



