DATA() { EXPOSED;

DATA EXPOSED SPECIAL

Around the Clock with Azure SQL and Azure Data Factory

Americas
February 3, 2021
09:00 - 17:00 PT

Asia
February 4, 2021
09:00 - 17:00 SGT

HOSTED BY
Wee Hyong Tok & Anna Hoffman

16 Sessions | 2 Ask the Expert Panels | 1 Hackathon

DATA() { EXPOSED;

Best practices using Azure SQL Database as Sink in ADF

Silvano Coriani – Program Manager, Azure SQL

Agenda

- Azure SQL Database connector
- Security
- Copy Activity
- Best practice for loading data into Azure SQL Database
- Scenarios
  - Append only
  - Upsert (insert new rows / update existing rows)
  - Custom bulk load logic
- Mapping data flows

Azure SQL Database connector

Connectivity “from and to” Azure SQL Database for:	Security	As Source	As Sink
• Copy activity • Mapping data flow • Lookup (retrieve a dataset to pass to subsequent activities) • GetMetadata	• Supports SQL authentication and Azure Active Directory • Service principal and Managed Identity	• Query • Stored procedure • Parallel copy	• Uses bulk load where possible • Create destination table if not exists • Invoke stored procedure with custom logic during copy

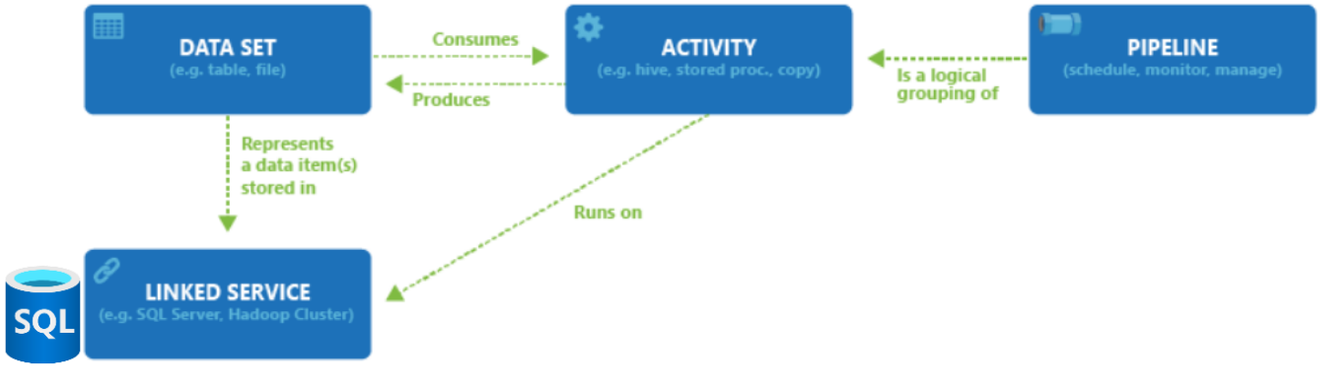

Copy Activity

- Copy Activity in Data Factory copies data from a source data store to a sink data store.

- Uses the Integration Runtime specified at the dataset level
  - Can be in Azure or Self-hosted
  - Can use a managed VNET and connect to Azure SQL through Private Link
- Azure SQL Database as source
  - Leverage partitioning and parallel copy
- Consider retry / retry interval for robust copy execution
  - Mandatory for Serverless tier
- Other data transformation activities can target Azure SQL Database as sink
  - Azure Function
  - Stored Procedure
  - Data Flow
  - Databricks Notebook

Best practice for loading data into Azure SQL Database

- Consider DIU and parallelism settings
  - Adjust (up to 32) to get required throughput/cost ratio

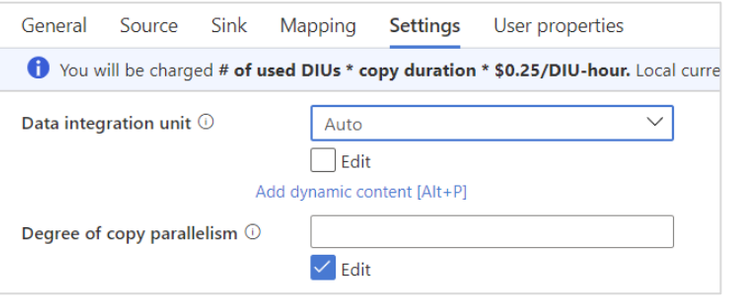

- Append data -> uses bulk insert
  - Can control attribute like batch size and max concurrent connections

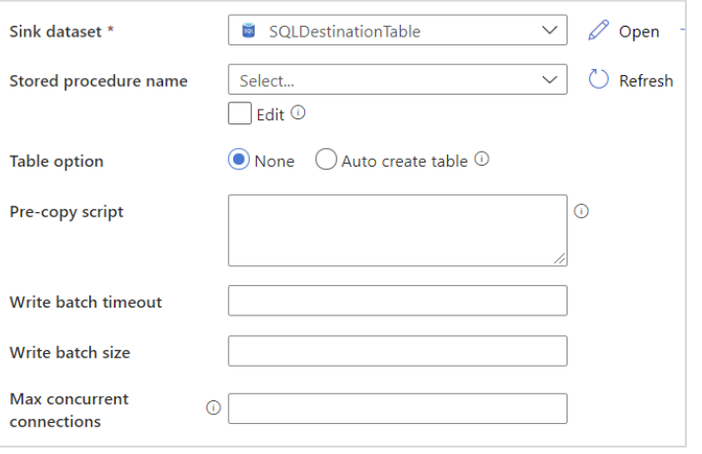

- Upsert -> bulk load in staging table and use MERGE or INSERT/UPDATE
  - Recommended for very large datasets
  - Consider a #temptable if writelog is an critical

Best practice for loading data into Azure SQL Database

- As an alternative, you define a Stored Procedure, invoked once for each batch and that receives a TVP containing the rows to insert
  - Recommended for mid-large datasets and for complex insert logic where you want to control all aspects of ingestion (e.g. control TABLOCK, etc.)

```
CREATE TYPE [dbo].[MarketingType] AS TABLE(
    [ProfileID] [varchar](256) NOT NULL,
    [State] [varchar](256) NOT NULL,
    [Category] [varchar](256) NOT NULL
)
```

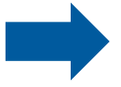

```
CREATE PROCEDURE spOverwriteMarketing @Marketing [dbo].[MarketingType] READONLY, @category varchar(256)
AS
BEGIN
MERGE [dbo].[Marketing] AS target
USING @Marketing AS source
ON (target.ProfileID = source.ProfileID and target.Category = @category)
WHEN MATCHED THEN
    UPDATE SET State = source.State
WHEN NOT MATCHED THEN
    INSERT (ProfileID, State, Category)
    VALUES (source.ProfileID, source.State, source.Category);
END
```

- Consider using proper Service Tier and compute size, check resource limits (e.g. max log rate). All data loading best practices and techniques (e.g. indexing, partitioning, etc.) are still relevant!

Max log rate (MBps)	24	48

Performance tuning tips:
Sink Azure SQL Database: The DTU utilization was high during the copy activity run. To achieve better performance, you are suggested to scale the database to a higher tier than the current 250 DTUs. Refer to this document.

```
SELECT primary_max_log_rate,primary_max_log_rate FROM sys.dm_user_db_resource_governance
```

Demo: optimize data loading in Azure SQL Database

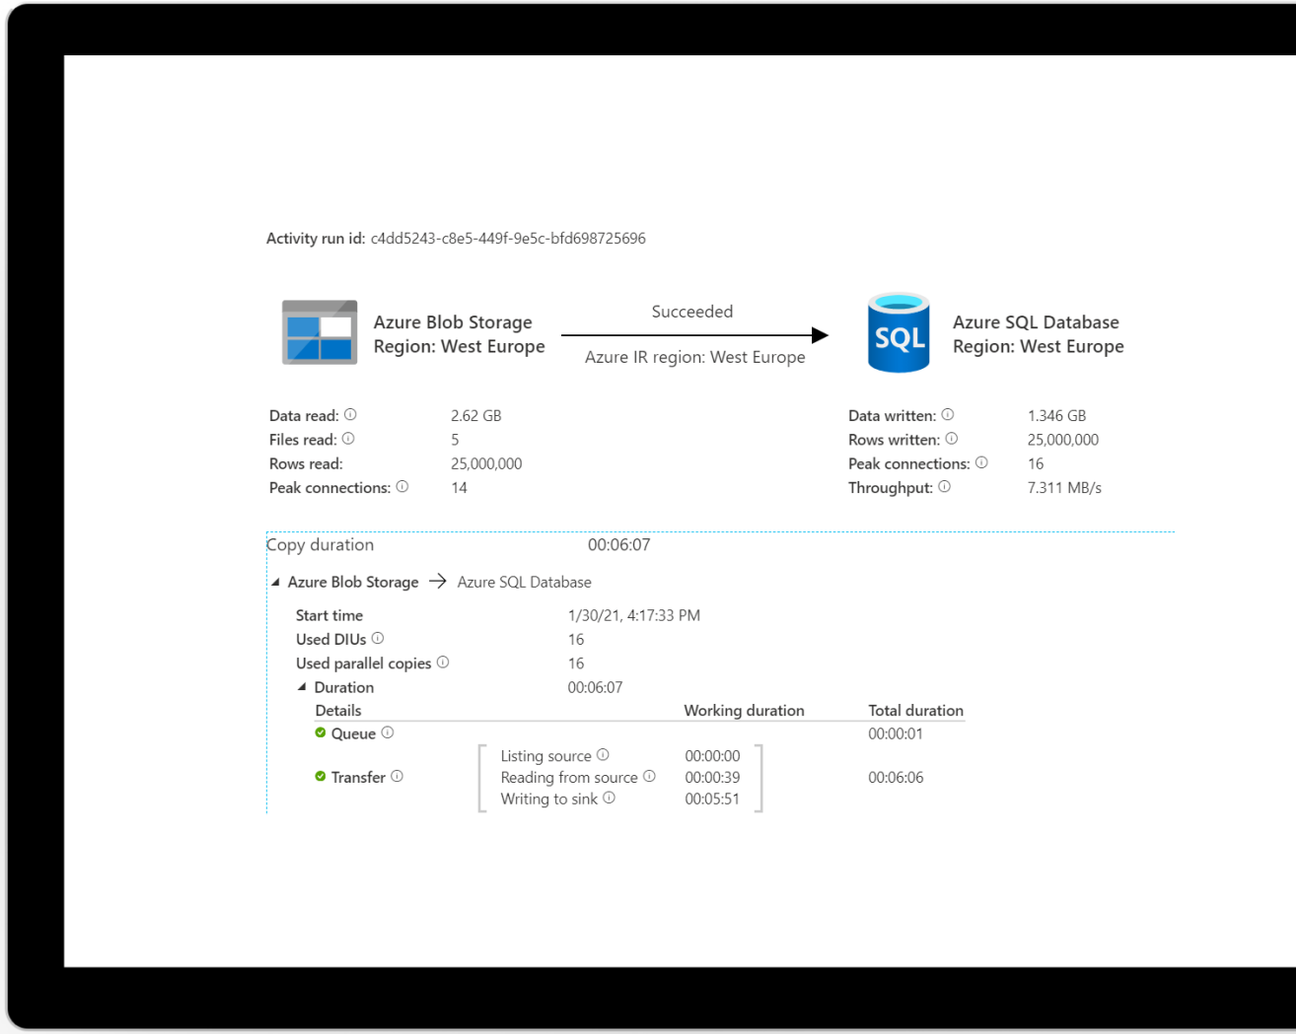

Mapping Data Flow

Visual data transformation engine based on scaled-out Apache Spark clusters

Recommended for large and complex data transformation tasks

Can read and write to Azure SQL Database tables

As source

- Define a read query (can use UDFs)
- Can define attributes like Batch size and Isolation level

As sink transformation

- Define what operations are allowed (insert, delete, upsert, update)
- Define a key column to determine which row to alter (support composite keys)
- Table action (recreate, truncate, none)
- Batch size (trade off between insert efficiency and memory usage)
- Use TempDB: select between using a global temp table or a persisted table

Error handling

- Control if fail or continue on error
- Determine transactional behavior (single transaction or batches)
- Output rejected data to Azure Blob Storage or ADLS Gen2

Recap

- Azure Data Factory is a great option for creating data movement and transformation solutions from and to Azure SQL Database.
- Copy Activity is indicated for most common data movement scenarios like Append, Upsert and custom logic.
- Understand Data Integration Unit and parallelism impact is critical to get maximum performance during copy operations.
- Prefer bulk insert into staging tables followed by T-SQL based transformations (ELT-like) to get best performance in append or upsert scenarios.
- If complex logic during inserts is still required, leverage invoking Stored Procedure and TVPs and evaluate proper batch size (depending on table size, optimal could be between 50k and 250k rows).
- Consider all usual Azure SQL Database best practices for schema and indexing optimizations (e.g. partitioning, clustered columnstore, etc.)
- Azure SQL Database capacity planning principles will help setting the stage

DATA() {
EXPOSED;

Learn with us!

View our on-demand playlist:
aka.ms/azuresqlandadf

@AzureSQL
@AzDataFactory

Microsoft Azure